

Reproducible Performance Evaluation Pipelines for Ceph

Mariette Soupe
UC Santa Cruz
msoupe@ucsc.edu

Abstract—Ceph, a free-software distributed system, is commonly studied and practiced among students with a background or interest in Computer Science. Setting up any distributed system is not a straightforward task and Ceph is no exception. This paper provides a reproducible automated setup of a Ceph distributed system with minimal configuring to first time users of Ceph or an experienced user on a macOS or Linux environment for end to end testing.

I. INTRODUCTION

A. Background

Setting up a distributed system like Ceph is hard and complex because there are many different components that make up Ceph. The goal of this project is to implement a workflow for Ceph experimentation. Instead of going the traditional route of manually installing everything in ad-hoc ways, a SciOps (DevOps for science) methodology is followed to create an experimentation workflow for Ceph. This results in having reproducible studies, transparency (open science), experiments that are amenable to collaboration and extension. In summary, we apply the SciOps methodology to Ceph experimentation that is typical of storage and data-management R&D settings (e.g think of companies such as Toshiba, WD, Samsung, etc.). This report describes the process of how the workflow was implemented and how it can be used.

The audience of this project is targeted towards several groups such as any newcomer of Ceph or existing users wanting a simple workflow for Ceph experimentation. Anecdotally, when students take a course in distributed systems and need to setup a distributed environment, it can take several weeks *just* to get the system up and running. If a student were in a ten week class, that does not give students enough time to play around with the system and report interesting observations or results since time is very limited. Reducing the overhead of the setup time will allow for more time working on a project versus spending time getting the system up and running. The workflow that has been implemented will allow for a variety of testing, but this paper focuses on one benchmarking test which will be performed as a proof of concept basis. Some of these tests which can be performed within Ceph include performance, scalability, correctness, availability, and overhead. The contributions of this project include:

- Applying SciOps methodology to implement a Ceph experimentation workflow;
- A template to deploy and test a Ceph cluster.

The remainder of this paper is organized as the following, Section II describes the approach and technologies used to create the workflow. Section III goes into detail about the different stages in the pipeline. Section IV reflects on the challenges while creating the experimentation workflow. Section V describes the experimental results and outcome of the experimentation workflow. Lastly, Section VI describes the future work and how this project can be further developed.

B. Ceph

Ceph [1] is an open source software that provides excellent performance, reliability, highly scalable objects, block, and file-based storage in a distributed system. Other features that are integrated within are having no single point of failure, uses commodity hardware and dynamically increase or decrease nodes where the system will be able to fully recover with underlying algorithms.

As previously mentioned Ceph is scalable. A Ceph cluster consists of many roles such as monitors, agents, OSDs, clients, managers, and rgws. For this project a small Ceph cluster is created, but the experiment is agnostic to the size of the cluster. When first starting out with Ceph, there is a minimum requirement of three nodes needed to get the system running. Each of these three nodes have their own role: a monitor node, one object storage device (OSD) node, and a client node. The monitor node contains a “cluster map” which is the set of maps comprising of the roles of a monitor, OSDs, placement group (PG), MDS, and CRUSH map [2]. These five different maps in the “cluster map” have knowledge of the whole cluster topology. Moreover, the monitor node provides logging and authentication services which are used when setting up the system. The OSD node is in charge of storing objects on a local file system and providing access to them over the network [2]. It is also important that there is one monitor node in a Ceph cluster because the monitor node acts as a leader for the cluster. Additionally, it is important to always have an odd amount of nodes in the system because we want the system to reach consistency so if there is an even amount of nodes there is a possibility of not obtaining a majority. When there is no majority, a system can get stuck and not reach consistency. For a proof of concept this paper focuses on a small system.

II. APPROACH

In order to get a Ceph cluster running for experimentation, there are different tools that are used to automate and create

a Ceph system; Cloudlab, Ceph Benchmarking Tool, Ceph-Ansible, Docker, and Popper. Some of these tools will need user credentials meaning that a user will need to create an account which are Cloudlab, Docker, and Github. Popper does not need credentials and can be installed through the terminal via the command, `pip install popper`. The first step in the workflow is changing/writing code and then building that code. Next, resources get allocated via Cloudlab. Proceeding this step, Ceph-ansible is used to deploy the code and then the Ceph Benchmarking Tool is used to run the different tests to be conducted on Ceph. Lastly, Jupyter is used to help analyze the results.

Cloudlab [3] is an online service that hosts bare metal machines for users to create and compute a cloud environment. This is where the resources come from. It is also important to note that since Cloudlab is a hosting service, requested machines are only available for a certain amount of time. Data that has not been saved before the expiration time will be erased from the machines. Since the allocation of the machines, also referred as nodes throughout the paper, is automated it doesn't matter too much that the machines expire because the machines can get reallocated fairly quickly.

Docker [4], a container platform, helps containerize packages and environments which can be shared among multiple applications. Docker will help create the Ceph environment into a container so all of Ceph's dependencies will be encapsulated together. As mentioned in the abstract, this workflow is only for MAC and Linux operating systems because of Docker. Docker does not work as easily and nicely on a Windows operating system which was observed from past experience. Once Docker has been installed onto a user's machine, then the obtaining images on Docker will be easy to pull from Docker hub which consists of many images. For this workflow, a base Docker images will contain all of the dependencies for some of the other mentioned tools will be contained within a container.

Popper [5] is a CLI tool and convention to create reproducible scientific articles and experiments. The convention is based on the open source software (OSS) development model which creates self-contained experiments that doesn't require external dependencies than what is already contained in the experiment. The Popper convention uses a pipeline that consists of shell scripts which executes an entire experiment. When a pipeline has been initialized, the pipeline consists of several default stages: the setup, run, postrun, validate, and teardown stage. These default stages are not required for every experiment and stages can be renamed accordingly to what makes sense for your experiment. For this experiment the stage names that are used are setup, deploy, run-benchmarks, teardown, and validate. In the setup stage, a user would usually download all the necessary files to run the project. These files are, for example, data files, libraries, and other dependencies. The run stage executes the script that is used to run the main part of the experiment. The postrun stage is where data can be manipulated so it is setup nicely to the validation stage. The validate stage is where a user would display the results obtained in the postrun stage. This stage could be used to open a log

file that shows the results of the experiment or run a script that graphs and displays the results, for our case Jupyter will be used to showcase the results.

Ceph-ansible [6] is an Ansible playbook that helps automate the installation of Ceph to create a Ceph cluster. Furthermore, Ansible is a software provisioning service which helps prepares an inventory file that defines a host. With this inventory file and host information, hosts roles are also defined to create a Ceph cluster.

Ceph Benchmarking Tool [7] is a tool that is used to automate different testing harnesses on top of a Ceph cluster for performance and benchmarking. There are four main benchmarking modules this tool provides; radosbench, librbd fio, kvmrbd fio, and rbd fio. CBT is a layer on top of Ceph and doesn't install Ceph packages to create a cluster which is why Ceph-ansible is used to create the cluster.

Jupyter [8], also commonly known as Jupyter Notebook, is an open source software interactive tool that integrates visuals, narrative text, and other forms of media in a single document. Jupyter is used in this workflow to help visualize our results which can then be analyzed for further testing.

III. PIPELINE

A. Prerequisites

There are four main prerequisites to run this experiment; Popper (v1.1.2), Docker (v2.0.0.3), a Cloudlab account and Github account. Any other additional dependencies that are needed to set up the cluster and benchmarking tool are contained in a Docker container, therefore no additional installation are required.

B. Workflow

Lst. 1 shows the complete pipeline for this experiment workflow. This pipeline consists of the following main stages; setup, deploy, run-benchmarks, teardown, and validate. The main goal of this project is not only to create a Ceph experimentation environment but to also create that environment with minimal configurations. After a user has the prerequisites installed on their machine, a user can run the Ceph experimentation environment, with the current default cluster configuration and benchmarking test, with two commands:

```
cd Ceph/pipelines/ceph
popper run
```

In addition to the scripts to conduct the experiment, there are also other folders in the pipeline such as CBT, Ceph-ansible, Docker, and GENI. These folders indicate the different tools are being used throughout the experiment. The use of each tool will be further explained in their appropriate stages.

C. Setup.sh

The setup stage consists of requesting resources. As mentioned earlier, the resources for this experiment will come from Cloudlab. For the resources that are being used for this project, Ceph is using Cloudlab's cluster Clemson and type c6320.

Listing 1 Contents of Ceph Popper repository.

```

Ceph-repo
| README.md
| .popper.yml
| pipelines
|   |-- ceph
|   |   |-- setup.sh
|   |   |-- deploy.sh
|   |   |-- run-benchmarks.sh
|   |   |-- teardown.sh
|   |   |-- validate.sh
|   |   |-- cbt/
|   |       |-- ceph.client.admin.key
|   |       |-- ceph.conf
|   |       |-- rbdmap
|   |       |-- conf.yml
|   |   |-- ceph-ansible/
|   |       |-- group_vars/
|   |           |-- all.yml
|   |           |-- osds.yml
|   |       |-- purge_cluster.yml
|   |       |-- site.yml
|   |-- Docker/
|   |   |-- cbt/
|   |       |-- Dockerfile
|   |       |-- install.sh
|   |   |-- ceph/
|   |       |-- Dockerfile
|   |       |-- install.sh
|   |-- geni/
|   |   |-- clemson.xml
|   |   |-- machines
|   |   |-- release.py
|   |   |-- request.py
|   |   |-- renew.py
|   |   |-- monitor_config.py
| paper
|   |-- build.sh
|   |-- figures/
|   |-- paper.md
|   |-- paper.pdf
|   -- references.bib

```

Lst. 2 shows the Clemson cluster hardware specifications for the type c6320. The Clemson cluster was chosen versus the other available clusters because of the amount of storage this specific hardware has.

For this experiment, using Cloudlab, there are five environment variables that need to be declared in order to begin the request of the nodes. The following environmental variables are; CLOUDLAB_USER, CLOUDLAB_PASSWORD, CLOUDLAB_PROJECT, CLOUDLAB_PUBKEY_PATH, CLOUDLAB_CERT_PATH.

Listing 2 Cloudlab Clemson c6320 hardware specifications.

CPU	Two Intel E5-2683 v3 14-core CPUs at 2.00 GHz (Haswell)
RAM	256GB ECC Memory
Disk	Two 1 TB 7.2K RPM 3G SATA HDDs
NIC	Dual-port Intel 10Gbe NIC (X520)
NIC	Qlogic QLE 7340 40 Gb/s Infiniband HCA (PCIe v3.0, 8 lanes)

Once the following environmental variables have been declared, these variables are passed into a Docker container to request resources autonomously through the GENI tool. GENI is an API that allows for ... Furthermore, the request API allows users choose which hardware they want to use. For our case Clemson c6320 is chosen. After the request of the resources completes, an `cl-clemson.xml` file is formed with all of the specifications for X amount of nodes, where X is the amount of nodes a user wants for their experiment. Next, a `machines` file is written to group the allocated resources from Cloudlab. Lastly, the monitor IP address gets copied from the `cl-clemson.xml` into the `all.yml` file in the `ceph-ansible` folder to prepare for the deployment of the cluster. Below shows which files are used for this stage and the descriptions of a files' functionality. Lst. 3 shows a table of the files and descriptions of their functionality. Lst. 3 shows the files that are used in the pipeline to allocate resources from Cloudlab.

Listing 3 Setup stage file breakdown

```

clemson.xml
- Specifications for each node
machines
- Allocated resources grouped
request.py
- Request resources from Cloudlab
renew.py
- Renew resources from Cloudlab
monitor_config.py
- Copying monitor IP address for
  deploy stage

```

D. Deploy.sh

After allocating resources, a Ceph cluster can be deployed. The tool that is used to create the Ceph cluster is Ceph-ansible. Ceph-ansible requires a few configuration files to deploy the cluster which can be referred to Lst. 1. First, all of the allocated nodes gets cleared and erased to make sure no previous version of Ceph has been installed. This will guarantee a clean slate when Ceph is installed. Next, the cluster design gets deployed and the monitor and OSDs roles are assigned to server groups. Lst. 4 shows the files that are used in the pipeline to create and configure a Ceph cluster.

Listing 4 Deploy stage file breakdown

```
all.yml
- Cluster network configuration settings
osd.yml
- Osd configuration
site.yml
- Defined deployment design and assigns
  role to server groups
purge-cluster.yml
- Purge and clean out any existing Ceph
  Installation on the nodes
Dockerfile
- Container containing Ceph-ansible
install.sh
- Clone Ceph-ansible and requirements
```

The way to verify that the Ceph cluster is up, is by checking the status of the Ceph cluster by one of three ways. First, looking at the end of the log from the deploy stage or secondly ssh into the monitor node and use the command `sudo ceph -s` to view the status of the cluster. Lst. 5 shows a sample output of a running cluster. Note that `health` has a warning. That warning can be ignored because in the configuration files when deploying the cluster, the `mgr` role was omitted. Another way to make sure that the cluster is configured properly is verifying that the OSDs are up and running. This can be checked by looking at the `services > OSD` from the output. In this case, it is shown that two OSDs are up. If the OSDs are not up and running, that indicates a cluster doesn't exist.

Listing 5 Cluster configuration output.

```
cluster:
  id:      xxx-xxx-xxx-xxx-xxxxx
  health:  HEALTH_WARN
           no active mgr

services:
  mon: 1 daemons, quorum node0
  mgr: no daemons active
  osd: 2 osds: 2 up, 2 in

data:
  pools:   0 pools, 0 pgs
  objects: 0 objects, 0 B
  usage:   0 B used, 0 B / 0 B avail
  pgs:
```

E. Run-benchmarks.sh

Once a Ceph cluster is up and running, tests can be conducted on the cluster. As mentioned previously there are

four different modules that CBT, Ceph Benchmarking Tool, provides; `radosbench`, `librbd fio`, `kvmrbd fio`, and `rbd fio`. Of these four testing modules, this pipeline conducts the `radosbench` test. Lst. 6 shows the files that are used in the pipeline to run the Ceph benchmarking tool.

Listing 6 Run-benchmark stage file breakdown

```
ceph.client.admin.key
- Key for the client.admin user
ceph.conf
- Ceph configuration settings
conf.yml
- Test configuration files to create
  parametric sweeps of tests
Docker-file
- Container containing cbt
install.sh
- Clone cbt and requirements
```

F. Teardown.sh

Once a user has completed their experiment there are two ways to terminate all of the allocated machines back to Cloudlab which will erase all of the data that was installed. The first method is letting the resources expire. The second way is using Cloudlab's release API to release the resources back into the resource pool. The teardown stage script will differ depending where the resources have been allocated, just like the setup stage's script.

G. Validate.sh

Lastly, in the validate stage, the data from the tests that have been conducted will be analyzed using Jupyter Notebook which is the last stage of the workflow.

IV. CHALLENGES

With any project, there are always obstacles that are faced through the process. The first main challenge going into this project was having minimal knowledge of Ceph. Having minimal knowledge of Ceph required me to catch up on the understanding of Ceph and not having prior experience creating a Ceph cluster. So while learning about how this environment workflow was supposed to set up, I had to also gain some understanding of what the environment does and how it works. This brings up the next challenge while working on this project and the one of the main purposes of this project, creating a Ceph cluster.

A key part of this workflow is for the pipeline to be reproducible. Incorporating reproducibility from the beginning is important so this entire workflow can be portable from one machine to another. Being a newcomer to Ceph and automating the whole process while keeping reproducibility in mind from the start was difficult because there was a lot of different components to keep track from the very beginning. Despite that difficulty, creating a reproducible workflow from the beginning

helped when I encountered issues and needed assistance for debugging.

Another challenge that pushed back the timeline of getting this pipeline working was having to change cloud hosts two times along the way. The first cloud host subscription expired, so then another alternative was suggested and create a local Ceph system. However, the local distribution failed due to permission rights on a Linux - Ubuntu machine and having to change some hardware settings that I didn't want played with. In hindsight a virtual machine would have been a better solution for the local Ceph system, but another main goal of the project was not to use a virtual machine. Virtual machines are useful, but require more installation steps versus just using Docker containers.

Another one of the challenge encountered was allocating nodes in Cloudlab. Cloudlab nodes consists of a variety of hardware, so when first starting out and allocating nodes, the nodes that were being allocated from the Cloudlab Utah cluster where the node had 256GB on the disk. There was a memory storage limit that was encountered since Ceph has a lot dependencies. Switching to the Cloudlab Clemson solved this issue.

V. RESULTS

VI. FUTURE WORK

Thus far this project has provided a template on how to get started on creating and conducting a test on a Ceph cluster with a default test harness. One feature that can be considered is incorporating other resource facilities for node allocation. As of right now, Cloudlab is the main facility this project is gathering its resources from which is not accessible for everyone. Cloudlab is for academia and is not open to everyone compared to Amazon Web Services, as an example. Adding other resource facilities such Amazon Web Services (AWS) or Chameleon will allow for a greater audience to use this experimentation workflow. Furthermore, using other resource facilities will not change the Ceph experimentation environment. The only change that'd occur are the setup and teardown stages of the pipeline because the main middle stages only require an inventory of the hosts and its roles for creating and testing in the cluster.

VII. CONCLUSION

In conclusion, building a Ceph distributed system is not an easy, straight forward task because there are different components that need to be put together. Additionally, it can take multiple weeks just to get the system running so in a class setting where weeks are limited most of the time is spent troubleshooting and debugging the system versus spending time obtaining interesting results to analyze and then make further tests within Ceph. This workflow, once understood, provides a lightweight automated, reproducible Ceph experiment environment.

REFERENCES

- [1] S.A. Weil, S.A. Brandt, E.L. Miller, D.D. Long, and C. Maltzahn, "Ceph: A scalable, high-performance distributed file system," *Proceedings of the 7th symposium on operating systems design and implementation*, USENIX Association, 2006, pp. 307–320.
- [2] "Ceph docs," <http://docs.ceph.com>.
- [3] "Cloudlab."
- [4] "Docker."
- [5] I. Jimenez, M. Sevilla, N. Watkins, C. Maltzahn, J. Lofstead, K. Mohror, A. Arpaci-Dusseau, and R. Arpaci-Dusseau, "The popper convention: Making reproducible systems evaluation practical," *Parallel and distributed processing symposium workshops (ipdpsw)*, 2017 *ieee international*, IEEE, 2017, pp. 1561–1570.
- [6] "Ceph ansible."
- [7] "Ceph benchmarking tool."
- [8] "Jupyter notebook."